

Deep Learning – Assignment 2

Face Verification with Siamese / Metric-Learning Networks

Amit Ner-Gaon (211649801) · Chen Frydman (208009845)
Ben-Gurion University of the Negev - Faculty of Computer & Information Science

1. Compute budget and fairness constraints

Stated before any results, as required by the assignment.

- **Hardware:** one NVIDIA RTX 4090 GPU, 48 GB system memory, 12-hour interactive session (`sinteractive --gpu rtx_4090:1 --time 12:00:00 --mem 48`). PyTorch, single-GPU, fp32.
- **Budget metric:** All configurations in Experiments 1 and 2 were trained using the same maximum epoch limit (EPOCHS = 60) and the same early stopping criterion (patience of 15 epochs without improvement in validation AUC). These values were chosen based on the findings of *“An Empirical Study on the Correlation between Early Stopping Patience and Epochs in Deep Learning”* by Bootan M. Hussein et al., which evaluated different patience values for CNN training. A relatively high epoch ceiling was selected due to the complexity of the face verification task. Early stopping prevents fast configs from being penalised by an unnecessarily long ceiling and prevents slow configs from reading as 'diverging' when they were just under-trained.
- **Hyperparameter search budget:** K = 4 trials per configuration, spent entirely on the margin (we did not also sweep the learning rate, we fixed Adam at $\text{lr} = 1\text{e-}3$ across all configs).
- **Input pipeline & augmentation:** identical resize, normalization, and the affine-distortion policy from Koch et al. §3.2 for every configuration within an experiment.
- **Evaluation episodes:** a single deterministic set of N-way one-shot episodes (per N) is generated once and reused across every model.
- **Seeds:** SEED = 42 for headline single-run numbers, SEEDS = (0, 1, 2) for the multi-seed re-runs of the best model per experiment.

2. Dataset (LFW-a)

We use the aligned LFW-a variant with the train/test pair splits supplied by `pairsDevTrain.txt` and `pairsDevTest.txt`. The test split contains identities not seen during training (verified: train/test identity overlap = 0). **Identities and pair counts:** 2,132 train identities, 963 test identities, 1,100/1,100 train positive/negative pairs and 500/500 test positive/negative pairs.

Images per identity (train): min = 1, mean = 3.14, median = 1, max = 530 - a heavy-tailed distribution dominated by single-image identities (histogram in Appendix Fig. A1).

Validation split. We carve ~10% of training identities (not pairs) into a held-out validation fold so identity-disjointness with train is preserved. Naively filtering the existing negative-pair list with this rule retains only $\sim \text{val_frac}^2 \approx 1\%$ of negatives - we observed 13 negative val pairs at $\text{val_frac} = 0.1$, which makes val-AUC dominated by sampling noise. We therefore **regenerate** validation negatives by sampling pairs of distinct val identities until the count matches val positives. Final split: train +978 / -887, val +122 / -122, train/val identity overlap = 0.

3. Methodology and experimental process

This section documents the experimental journey the grader explicitly asked for: the hypotheses tested, how parameter values were chosen, and what we learned along the way.

Hypotheses pre-registered before running the experiments.

- **H1 (loss):** triplet loss with semi-hard mining outperforms contrastive, which in turn outperforms Koch-style BCE on L1 distance.
- **H2 (mining):** semi-hard mining outperforms random triplet sampling at matched compute (this is the FaceNet finding).
- **H3 (backbone):** ResNet-18 (residual) modestly outperforms a parameter-matched Koch CNN (plain stacking).
- **H4 (pretraining floor):** frozen ImageNet ResNet-18 with cosine similarity is competitive with the Koch-paper baseline but loses to trained metric-learning models.

How parameters were chosen. With $K = 4$ trials per configuration available, we spent the budget on the loss margins, which we judged (and the literature confirms) have the largest single effect on training stability for metric losses. Adam was fixed at $\text{lr} = 1\text{e-}3$ with no schedule. Margin grids: contrastive {0.5, 1.0, 1.5, 2.0} on L2-normalised embeddings (effective range [0, 2]), triplet {0.1, 0.2, 0.3, 0.5} around the FaceNet default $\alpha = 0.2$. Each trial ran for $\text{EPOCHS}/3 = 20$ (short-schedule sweep), and the margin with the best validation verification accuracy was retrained for the full budget. The sweep chose **contrastive margin 0.5** and **triplet margin 0.2**. We did not sweep the learning rate (would have exceeded $K = 4$ per loss) and we did not sweep augmentation strength (the augmentation policy is held constant across configurations as a fairness constraint).

Methodological iterations made during this project. Three discoveries during early runs changed the methodology:

- **Semi-hard mining selection.** A first implementation picked the negative with the largest d_{an} within the semi-hard window (the easiest semi-hard), producing near-zero loss per surviving triplet and letting random sampling win. The correct choice is $\text{argmin}(d_{an})$ within the window - the hardest semi-hard. After the fix, semi-hard and random sit within noise of each other (see Experiment 1).
- **Validation negatives.** The original identity-filter retained only 13 val negatives, val-AUC became too noisy to use for model selection or threshold choice. We rebalanced by regenerating val negatives within val identities. After the fix, the val→test AUC gap for the best model dropped from ~ 0.11 to ~ 0.01 .
- **Early-stop patience.** Initially 10. Triplet runs with strict-skip semi-hard mining have a sparse-gradient warm-up phase that can stall val-AUC for several epochs, we raised patience to 15 to avoid stopping before the in-batch semi-hard window populates.

Threshold selection. Verification accuracy is reported at the threshold that maximises Youden's J statistic (TPR - FPR) on the validation ROC, the test set is never used to pick the threshold. Same protocol for every model (Koch BCE uses P(same) from the head, contrastive and triplet use negative L2 distance on L2-normalised embeddings, frozen ImageNet uses cosine similarity).

One-shot protocol. 250 episodes per $N \in \{2, 5, 20\}$, built once with SEED = 42 from test identities only and shared across every model. Each episode is (query, [N supports], correct_idx), the model is correct iff the same-identity support has the highest similarity to the query.

Tools and instrumentation. Parameter counts and MACs are reported via the thop library (`pip install thop`, pytorch-OpCounter, v0.1) on a dummy $1 \times 1 \times 105 \times 105$ tensor for the Koch CNN and $1 \times 3 \times 112 \times 112$ for ResNet-18 - matching each backbone's training-time input. We report MACs directly and FLOPs = $2 \times \text{MACs}$ (explicit conversion). t-SNE via sklearn.manifold.TSNE (perplexity = 15, PCA init). ROC/AUC via sklearn.metrics. No metric-learning repositories were copied.

4. Experiment 1 - Loss function

Backbone held fixed (slimmed Koch CNN, EMBED_DIM = 128). Three losses are compared, with random-triplet sampling reported as an explicit ablation against semi-hard mining.

How each loss shapes the embedding. **Koch BCE on L1** learns a scalar similarity directly through a Linear(EMBED_DIM, 1) head on the elementwise $|z_a - z_b|$, the embedding space has no native metric, only the post-hoc weighted-L1 head ranks pairs. **Contrastive** induces a Euclidean geometry: same-identity points pulled toward distance 0, different-identity points pushed at least to the margin m . **Triplet** enforces a relative constraint $d(a, p) + \alpha \leq d(a, n)$ without pinning an absolute scale, which is why we L2-normalise embeddings before computing distances. **Role**

of margin. For contrastive, m sets the absolute scale of inter-class separation, too small and the loss collapses, too large and the loss can never reach zero. For triplet, α sets the minimum *relative* separation required between a positive and a negative, with L2-normalised embeddings ($\max d^2 = 4$), the FaceNet default $\alpha = 0.2$ was chosen by our $K = 4$ margin sweep.

Why mining matters. Random triplets are mostly trivial after a few epochs (most negatives are already at $d > d(a, p) + \alpha$ and contribute zero gradient), so the model stops learning, mining selects informative triplets that still violate the margin.

Model	Acc	AUC	N=2	N=5	N=20	Wall (s)
Koch BCE	0.647	0.686	0.724	0.392	0.104	140
Contrastive	0.704	0.761	0.768	0.432	0.284	136
Triplet (semi)	0.834	0.904	0.940	0.784	0.552	524
Triplet (rand)	0.813	0.914	0.920	0.768	0.524	587

Table 1: Experiment 1 results (single seed = 42). ROC overlay + train and validation loss curves in Appendix Figs. A2.

Findings (confirms H1, contradicts H2). Going from Koch BCE $\rightarrow$ Contrastive $\rightarrow$ Triplet adds +0.075 and +0.143 test AUC respectively - the monotonic loss ordering predicted by H1 is confirmed. The Koch BCE baseline barely beats chance on N=20 one-shot (0.104 vs random 0.05), confirming that a learned scalar similarity head produces embeddings whose pairwise structure is weak. Triplet (semi) achieves 0.904 AUC, triplet (rand) sits at 0.914 on this seed - random tied or slightly beat semi-hard, contradicting H2. The likely cause is our strict-skip semi-hard implementation: for many anchors the window ($d_{ap}, d_{ap} + \alpha$) is empty after a few epochs, so those triplets contribute no gradient, random sampling occasionally lands on a true hard negative and gets a full-strength update. The 3-seed mean for semi-hard is 0.894 ± 0.008 (Section 8), comfortably covering the random-triplet single-seed result. We interpret this as: mining is methodologically correct (and what FaceNet uses), but with strict-skip on a small dataset, the gradient sparsity neutralises the expected advantage.

5. Experiment 2 - Backbone

Loss fixed to triplet w/ semi-hard mining - the Exp 1 winner among the three required losses. Random-triplet sampling, although marginally ahead on single-seed test AUC, is explicitly disqualified by the assignment as 'not an acceptable triplet baseline' (only permitted as an ablation), so it is not eligible to fix the loss for Exp 2. Backbones compared at matched parameter count.

Backbone	Params	MACs	FLOPs (=2×MACs)	Wall (s)
Koch CNN	10,764,224	8.80×10^8	1.76×10^9	524

(slimmed)				
ResNet-18 (scratch)	11,242,176	4.87×10^8	0.97×10^9	702

Table 2: Parameter / MACs / FLOPs (thop on $1 \times 1 \times 105 \times 105$ for Koch, $1 \times 3 \times 112 \times 112$ for ResNet). Parameter difference: 4.3%, well inside the $\pm 20\%$ matching constraint. Koch has higher MACs per forward pass because of its larger 105×105 input and 10×10 first kernel, ResNet has more depth but smaller spatial footprint at 112×112 .

Backbone (loss = triplet semi-hard)	Acc	AUC	N=2	N=5	N=20
Koch CNN	0.834	0.904	0.940	0.784	0.552
ResNet-18 (scratch)	0.810	0.895	0.932	0.756	0.508

Table 3: Experiment 2 results (single seed = 42). Loss curves + ROC in Appendix Fig. A3.

Findings (disagree with H3). At matched parameter count and matched loss, Koch CNN matches or modestly beats ResNet-18 from scratch on every metric. Residual connectivity did NOT buy a measurable advantage on LFW-a at this scale. Plausible explanation: $\sim 1,900$ training identities with a median of 1 image each does not provide enough data to exploit the inductive bias residual blocks offer, the heavier first-conv stack in Koch happens to capture enough of the face-discriminative low-level structure. The result was robust across seeds (Section 8): Koch+triplet 0.894 ± 0.008 AUC vs ResNet+triplet 0.881 ± 0.005 AUC ($a > 1\sigma$ separation), and ResNet showed larger seed variance on N-way one-shot (e.g., $N=20 \pm 0.043$ vs Koch's ± 0.015). We would have predicted from the lectures that residual blocks ease optimisation on deeper networks - this is true in principle, but at 18 layers and small data, the optimisation advantage is not the bottleneck.

6. Experiment 3 - Frozen pretrained baseline

Frozen ImageNet-pretrained ResNet-18 (final FC replaced with Identity, all parameters frozen, no fine-tuning). Embeddings are the raw 512-D backbone features, pairs scored by cosine similarity, same test pairs and one-shot episodes as Experiments 1 and 2.

One caveat is that LFW-a is provided as single-channel images and ImageNet's ResNet-18 expects 3-channel RGB normalised with ImageNet mean/std. We replicate the grayscale channel three times and apply ImageNet normalisation, which produces a domain mismatch (ImageNet was trained on RGB photos) that likely lowers the frozen-baseline ceiling. Resampling to 112×112 (not 224×224) further moves it off the ImageNet distribution. Test AUC = **0.763**, verification accuracy = 0.684, N-way one-shot $\{2, 5, 20\} = \{0.752, 0.540, 0.312\}$.

Finding (confirms H4). The frozen baseline sits between the Koch-paper baseline (AUC 0.686) and every trained metric-learning model ($\text{AUC} \geq 0.895$). Pretrained ImageNet features are not competitive with trained metric learning on faces - but they are a strong sanity floor that the literal Koch reproduction fails to clear.

Koch BCE on L1 distance (AUC 0.686) does not beat the frozen ImageNet baseline (AUC 0.763), so the paper's original setup, evaluated faithfully at matched compute and dataset, is dominated by a zero-training feature extractor trained on an unrelated task.

No multi-seed re-run is reported for this baseline: with all weights frozen, no random augmentation, and a deterministic forward pass over a fixed test set, the embeddings (and therefore every reported metric) are bit-identical across seeds, so a 3-seed mean $\pm$ std would be trivially ± 0 .

7. Overall comparison and embedding analysis

Model	Acc	AUC	thr	N=2	N=5	N=20
Koch BCE	0.647	0.686	+0.499	0.724	0.392	0.104
Contrastive	0.704	0.761	-0.247	0.768	0.432	0.284
Triplet (semi-hard)	0.834	0.904	-1.134	0.940	0.784	0.552
Triplet (random ablation)	0.813	0.914	-0.874	0.920	0.768	0.524
ResNet-18 (scratch)	0.810	0.895	-0.949	0.932	0.756	0.508
ResNet-18 (frozen ImageNet)	0.684	0.763	+0.757	0.752	0.540	0.312

Table 4: Single-seed master table on the same test pairs and the same one-shot episodes. Single ROC overlay across all models in Appendix Fig. A4.

Embedding analysis for the best permitted model (selected by validation AUC over the assignment-permitted baselines: ResNet-18 (scratch), $\text{val_AUC} = 0.908$, $\text{test_AUC} = 0.895$). The val-AUC argmax across ALL single-seed candidates would have picked the random-triplet ablation ($\text{val_AUC} = 0.924$, $\text{test_AUC} = 0.914$), the assignment explicitly disqualifies random sampling as a triplet baseline (permitted only as an ablation), so we analyse the highest-val-AUC permitted model instead. This keeps section 5 and 7 consistent. Note that the 3-seed mean test AUC reported in §8 reverses this single-seed ranking: Koch+triplet (semi-hard) at 0.894 ± 0.008 beats ResNet-18+triplet at 0.881 ± 0.005 . We retain the single-seed val-AUC selection rule here to avoid using the test set in any selection step, the embedding analysis is therefore on the val-AUC winner, while §5/§8/§11 use the multi-seed-mean winner. At $\sim 1\sigma$ separation, the two embedding spaces are likely qualitatively similar (both are triplet semi-hard). 2D t-SNE of test-

set embeddings over 25 sampled test identities (up to 8 images each) in Appendix Fig. A5 shows tight, well-separated clusters for most identities, with some clusters smearing into neighbours (consistent with the failure cases in §9). Intra-class vs inter-class L2 distance on the same 25-identity sample: intra mean = 0.725 ± 0.295 ($n = 316$), inter mean = 1.355 ± 0.355 ($n = 7,812$). Inter-class mean is $\sim 1.9\times$ intra-class mean and the standard deviations are tight enough that the bulk of the distributions separate (histogram in Appendix Fig. A6), though the upper tail of intra and lower tail of inter overlap meaningfully - exactly the source of the residual verification errors ($1 - \text{AUC} \approx 0.10$ on test).

8. Multi-seed validation of the best model per experiment

Recipe	Acc	AUC	N=2	N=5	N=20
Koch + triplet (semi)	0.798 ± 0.007	0.894 ± 0.008	0.909 ± 0.010	0.751 ± 0.009	0.495 ± 0.015
ResNet-18 + triplet	0.790 ± 0.011	0.881 ± 0.005	0.895 ± 0.027	0.727 ± 0.007	0.439 ± 0.043

Table 5: Mean $\pm$ std across SEEDS = (0, 1, 2). Bar+error+per-seed-dot visualisation in Appendix Fig. A9. SEED = 42 (single-run headline) sat on the high side of the distribution for both recipes - the mean is the correct number to quote. AUC gap of 0.013 between backbones exceeds the combined $\sim\sigma$ of 0.009, so the architecture difference is weakly significant in our favour for Koch. Note that ResNet-18 has notably wider seed variance, particularly on N-way one-shot at large N.

9. Failure-case analysis (best model)

Six representative misclassifications of the best model on the test set, selected as the **most-confident** errors. Images in Appendix Fig. A7-A8.

- **False accept 1** (predicted same, different identities): two well-lit frontal portraits with similar hairlines, glasses, and skin tone - the model latches onto coarse appearance.
- **False accept 2**: heavy three-quarter pose on one image vs frontal on the other but similar facial hair and head-tilt - the small affine augmentation policy may not span this pose gap.
- **False accept 3**: similar age, gender, and ethnicity with neutral expression - hypothesised cause is the small dataset (median 1 image/identity) not letting the embedding learn fine identity cues.
- **False reject 1** (predicted different, same identity): one image in shadow, the other under bright stage lighting - illumination shift dominates the L2 distance.
- **False reject 2**: same person photographed years apart - age drift (skin texture, hair colour) is a known weak point for face verification without age-aware loss.

- **False reject 3:** occlusion (microphone / hand near face) on one image - the affine augmentation policy doesn't introduce occlusion, so the model has no trained invariance to it.

10. Threats to validity

What would change our conclusions. A longer LR schedule (cosine or step) could close the gap between Koch BCE and the metric losses, embedding dim 128 was chosen to match across backbones and could be raised to 256/512 - both choices would raise absolute AUC without necessarily changing the head-to-head ordering. Our semi-hard mining uses strict skip rather than softer alternatives (e.g., FaceNet's random-within-window or curriculum), which directly hurts H2 - under a different mining policy, semi-hard would likely beat random by a small but consistent margin.

What we did not control for. Per-subgroup performance (gender, ethnicity, age band) - LFW is known to be biased toward English-speaking, mostly-male, mostly-frontal celebrity faces.

What seed and budget choices may have hidden. Only 3 seeds for the best per-experiment model, the standard deviations are themselves imprecisely estimated. Random-triplet was reported single-seed, if multi-seeded its mean might land below semi-hard. Early stopping at patience = 15 epochs might still curtail a slow second-wind for triplet semi-hard - extending patience would be the next ablation. All thresholds are chosen on validation, never on test.

11. Conclusions

- **Loss matters more than backbone at this scale.** Koch BCE → Koch+triplet is +0.218 AUC, Koch+triplet → ResNet+triplet is -0.013 AUC across 3 seeds. The loss/sampling axis dominates by an order of magnitude.
- **Mining bought no measurable advantage under strict-skip semi-hard.** Random sampling tied with semi-hard. A softer mining policy is the likely fix, deferred as out-of-scope.
- **Architecture parity is essential.** At 10.76M vs 11.24M (4.3% gap), residual connectivity did not help on the mean and increased seed variance.
- **Frozen ImageNet is a strong sanity floor** that the literal Koch baseline does not clear. Trained metric learning still wins.

Appendix

Figures referenced from the body. All numbers in the body match the figures here exactly (extracted from the executed notebook).




Fig. A1 - Distribution of images per training identity. Heavy tail, median = 1.




Fig. A2 - Experiment 1: ROC overlay (all four losses) on the left, training loss (middle) and validation loss (right) for the same four runs. Different losses live on different scales - the shapes are comparable, the absolute magnitudes across rows are not.



Fig. A3 - Experiment 2: Koch+triplet vs ResNet-18+triplet. Same loss $\rightarrow$ magnitudes ARE comparable. ResNet's slower per-epoch progress is visible.




Fig. A4 - Single ROC overlay across every model in Experiments 1–3 (assignment requirement). The frozen ImageNet baseline clearly sits between Koch BCE and the trained metric-learning models.




Fig. A5 - 2D t-SNE of best-model test embeddings, 25 sampled identities $\times$ up to 8 images per identity. Legend (right) names each identity by its LFW folder name. Most identities form tight clusters, the few smeared clusters correspond to the failure cases.




Fig. A6 - Intra- vs inter-class L2 distance distributions (best model). Inter-class mean is well above intra-class mean, tail overlap is the source of test errors.

False accepts (same predicted, different identity)

score=-0.185



score=-0.195



score=-0.283



Fig. A7 - Top-3 most-confident false accepts (predicted same, actually different identities).

False rejects (different predicted, same identity)

score=-1.865



score=-1.685



score=-1.643



Fig. A8 - Top-3 most-confident false rejects (predicted different, actually same identity).




Fig. A9 - Multi-seed validation. One panel per metric, bar = mean across SEEDS = (0, 1, 2), error bar = ± 1 std, black dots = individual seeds, red x = the SEED = 42 headline value reported in the body. Koch+triplet beats ResNet-18+triplet on every metric and shows lower seed variance.